



LEADS 3-Day Camp

Session 2: Big Data Management: Relational Databases and SQL

Il-Yeol Song, Ph.D.
Professor
College of Computing & Informatics
Drexel University
Philadelphia, PA 19104
Song@drexel.edu
<http://www.cci.drexel.edu/faculty/song/>

Il-Yeol Song, Ph.D.

1



Hello!



Instructor


- Professor, College of Computing & Informatics, Drexel University,
- PhD in CS, LSU, Baton Rouge, USA, 1988
- PhD Program Director (2010-2015)
- CVDI Deputy Director (2012-2014)
- Research Topics: *Conceptual Modeling, Data Warehousing, Big Data Management and Data Analytics, Smart Aging*
- Elected as an **ER Fellow**, 2012
- Named an **ACM Distinguished Scientist** in 2013
- Received **Peter Chen Award in Conceptual Modeling** in 2015
- Four teaching awards from Drexel (1991, 2000, 2001, 2011) including **Lindback Distinguished Teaching Award** (2001)
- **Co-Editor-in-Chief**, Journal of Computing Science & Engineering
- **Consulting Editor**, Data & Knowledge Engineering Journal
- Co-Chair, **iSchool Model Data Science Curriculum Committee**
- **Chair, IEEE Big Data and Smart Computing (BigComp) Conference**
- Published about 200+ papers

Il-Yeol Song, Ph.D.



My Google Scholar Page

Google Scholar



Il-Yeol Song
[Drexel University](#)
 Verified email at drexel.edu - [Homepage](#)
[Conceptual modeling](#) [Data Warehouses](#) [Object-Oriented Modeling](#) [Big Data](#) [Smart Aging](#)

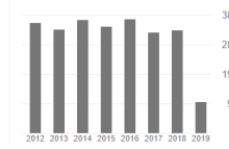
[FOLLOW](#) [GET MY OWN PROFILE](#)

TITLE	CITED BY	YEAR
A survey on ontology mapping N Choi, IY Song, H Han ACM Sigmod Record 35 (3), 34-41	682	2006
A UML profile for multidimensional modeling in data warehouses S Luján-Mora, J Trujillo, IY Song Data & Knowledge Engineering 59 (3), 725-769	349	2006
Analytics over large-scale multidimensional data: the big data revolution! A Cuzzocrea, IY Song, KC Davis Proceedings of the ACM 14th international workshop on Data Warehousing and ...	306	2011
Designing data warehouses with OO conceptual models J Trujillo, M Palomar, J Gomez, IY Song Computer 34 (12), 66-75	239	2001
Extending the UML for multidimensional modeling S Luján-Mora, J Trujillo, IY Song International Conference on the Unified Modeling Language, 290-304	135	2002
Integration of association rules and ontologies for semantic query expansion M Song, IY Song, X Hu, RB Allen Data & Knowledge Engineering 63 (1), 63-75	114	2007
XML-OLAP: a multidimensional analysis framework for XML warehouses BK Park, H Han, IY Song International Conference on Data Warehousing and Knowledge Discovery, 32-42	112	2005

Cited by

	All	Since 2014
Citations	4923	1834
h-index	34	20
i10-index	82	37

VIEW ALL





Il-Yeol Song, Ph.D.

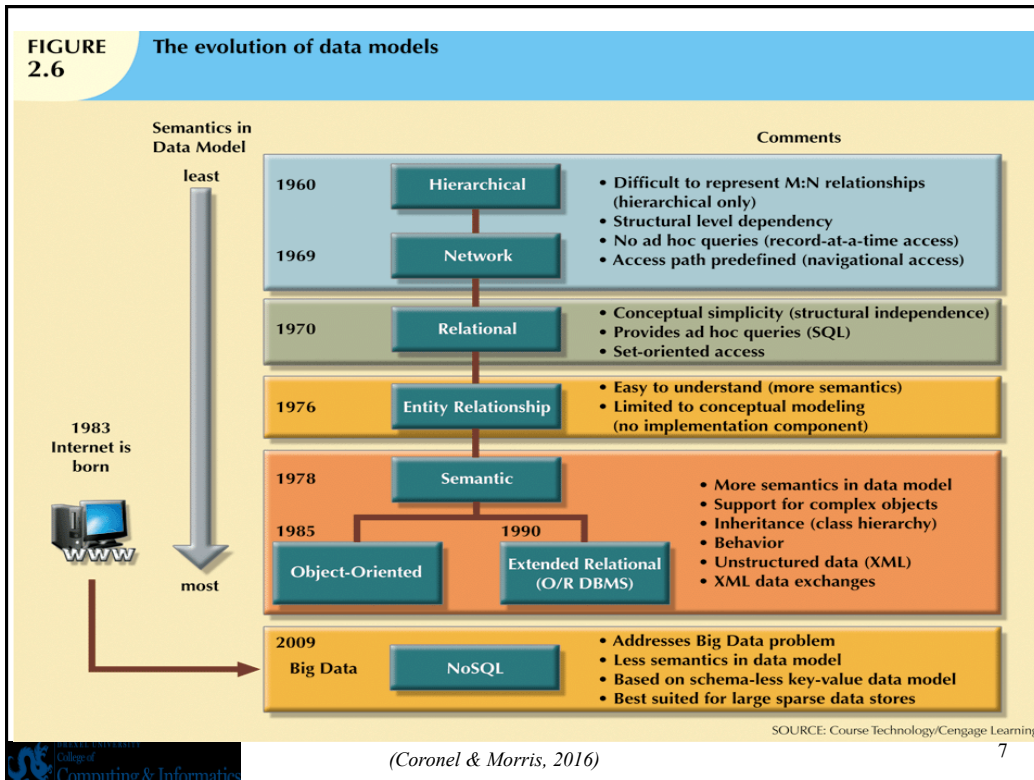
Quick Survey

- How many of you are familiar with:
 - Relational database concepts?
 - SQL?
 - Any experience of using relational databases (Oracle, MS SQL Server, MySQL, Postgres, DB2, etc)?
 - Entity-relationship Modeling?
 - Understand Difference between RDBs and NoSQL databases?
 - Used MongoDB?



Acknowledgement

- Materials from:
 - Drexel CCI INFO 605 and 606 Class (Il-Yeol Song)
 - Some diagrams and examples from
Carlos Coronel and Steven Morris, Database Systems: Design, Implementation, and Management, 12th Edition, 2016. Cengage Learning.



Basics in the Relational Model

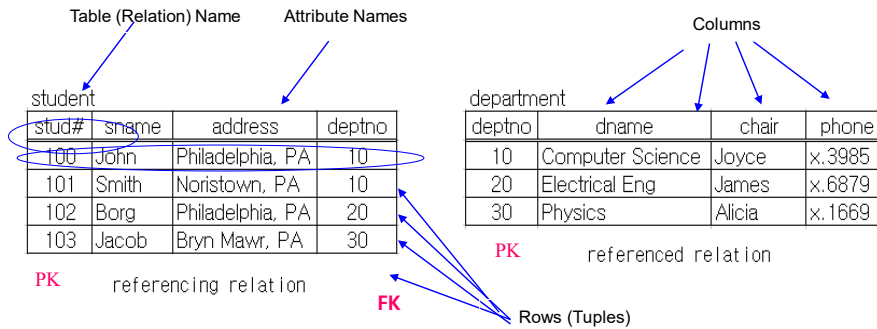
- Data representation: a set of tables with structured data
- Primary keys (Candidate key, Alternate keys)
- Foreign Key
- Referential Integrity constrains
- Integrity constraints in RDBMSs
- Null values
- Metadata and Data Dictionary
- Relational schema design using ER diagrams
- Normalization (FD, MVD, JD) and Normal Forms
- SQL (*Aggregation, Subquery, JOIN, PL/SQL* functions/procedures/triggers)
- ACID property in multi-user transactions
- Why relational databases have been popular in industry?

Terminology in RDB

Relational Schema

STUDENT (stud#, sname, address, deptno)

DEPARTMENT (deptno, dname, chair, phone)



A **Primary Key (PK)** uniquely identifies each row.

A **Foreign Key (FK)** is a set of attributes that are primary key values from another table.

Characteristics of Relational Model




TABLE 3.1 Characteristics of a Relational Table

1	A table is perceived as a two-dimensional structure composed of rows and columns.
2	Each table row (tuple) represents a single entity occurrence within the entity set.
3	Each table column represents an attribute, and each column has a distinct name.
4	Each row/column intersection represents a single data value.
5	All values in a column must conform to the same data format.
6	Each column has a specific range of values known as the attribute domain .
7	The order of the rows and columns is immaterial to the DBMS.
8	Each table must have an attribute or a combination of attributes that uniquely identifies each row.

Example of a DDL in SQL Defining a Relational Schema

stud#	sname	address	deptno
100	John	Philadelphia, PA	10
101	Smith	Norristown, PA	10
102	Borg	Philadelphia, PA	20
103	Jacob	Bryn Mawr, PA	30

deptno	dname	chair	phone
10	Computer Science	Joyce	x.3995
20	Electrical Eng	James	x.6879
30	Physics	Alicia	x.1669

referencing relation

referenced relation

```
CREATE TABLE Department (
    DeptNo      NUMBER(5)    PRIMARY KEY,
    DName       VARCHAR2(15) NOT NULL,
    Chair       VARCHAR2 (20),
    Phone       CHAR (10),
);
```


```
CREATE TABLE Student (
    Stud#       NUMBER(8),
    SName       VARCHAR2(20) NOT NULL,
    Address     VARCHAR2 (40),
    DeptNo      NUMBER (5),
    CONSTRAINT Stud_PK PRIMARY KEY (Stud#),
    CONSTRAINT Stud_FK FOREIGN KEY (DeptNo)
    REFERENCES Department (DeptNo)
);
```

ER Model: Popular Notations


Chen




Visio 's Relational notation




UML




IDEF1X



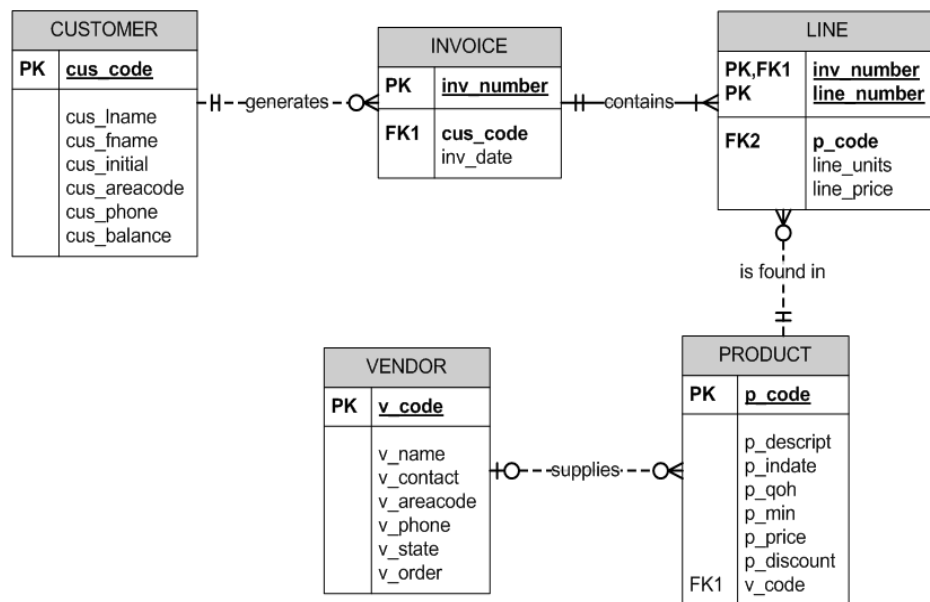
Crow's foot with Relational Notation in Visio



Visio IDEF1X with Crow's foot



ERD



(Coronel & Morris, 2016)

13


Why Relational Databases?

- Logical simplicity of the schema: Tables
- Easy, powerful, standard database language: SQL
- Transaction reliability: ACID property (Atomicity, Consistency, Isolation, Durability)
- Ad-hoc query processing
- Mature and reliable technologies
- Commercial investment for the last 40 years
- A large group of man-power and user bases

Il-Yeol Song, Ph.D.

14

Top Tools used in Data Science (2017)



7,955 responses

Source: <https://www.kaggle.com/surveys/2017>

Two major Components of SQL

Data Definition Language (DDL)

- CREATE TABLE
- ALTER TABLE
- DROP TABLE
- CREATE/DROP VIEW
- CREATE/DROP INDEX

Data Manipulation Language(DML)

- UPDATE TABLE
- INSERT INTO
- DELETE FROM
- **SELECT FROM**

DDL manages metadata
DML manages data

A Typical SQL Statement and Rules

STUDENT

ID	Name	EnrollYear

- **SELECT** enrollyear **FROM** student **WHERE** name = 'John';
- SQL statement comprises
 - **Keyword**
 - Case-insensitive
 - Basic command set: fewer than 100 English words
 - Space-independent
 - **User-defined word** (table names, attribute names)
 - Case-insensitive
 - **Data**
 - **Case-sensitive** (In MS Access, case-insensitive)
 - Ends with a semi-colon

An Overview of SQL Commands

DDL

--Changing attribute type

ALTER TABLE project

MODIFY (budget
NUMBER(9,2));

--Adding an attribute

ALTER TABLE project **ADD**
(manager CHAR(10));

-- Delete a column from a table

ALTER TABLE project
DROP COLUMN
manager;

--Removing a table from database

DROP TABLE project;

DML

--Inserting a row to a table

INSERT INTO project **VALUES** (1234,
'Perfect Project', NULL, 'John');

--Changing a value of attribute

UPDATE project **SET** budget =
1.1*budget **WHERE** projno > 1000;

--Deleting a row from a table

DELETE FROM project **WHERE**
manager = 'John';

SELECT Statement

SELECT	Specifies which columns are to appear in output
FROM	Specifies table(s) to be used
WHERE	Filters rows with conditions
GROUP BY	Forms groups of rows with same column value.
HAVING	Filters groups subject to some condition.
ORDER BY	Specifies the order of the output.

- *Only **SELECT** and **FROM** are mandatory.*
- *Order of the clauses cannot be changed.*

An Overview of Simple Commands

- | | |
|--|--|
| (1) SELECT * FROM
Instructor; | (6) SELECT * FROM Instructor
WHERE deptCode IS NULL ; |
| (2) SELECT DISTINCT fName
FROM Instructor; | (7) SELECT * FROM Instructor
WHERE deptCode IS NOT NULL ; |
| (3) SELECT fName, lName, ssn
FROM Instructor
WHERE deptCode = 'math'; | (8) SELECT * FROM Instructor
WHERE bonus BETWEEN 500 AND 1000; |
| (4) SELECT fName, lName
FROM Instructor
WHERE bonus > 1000; | (9) SELECT * FROM HR
WHERE hireDate
BETWEEN '01-MAY-2012' AND '31-MAY-2012'; |
| (5) SELECT * FROM Instructor
WHERE (bonus >= 500 AND bonus <= 1000); | (10) SELECT * FROM Instructor
WHERE bonus IN (100, 200, 300); |

An Overview of Simple Commands

- (11) SELECT * FROM
Instructor
WHERE position = 'assistant';
- (12) SELECT * FROM R
WHERE lName > 'S';
- (13) SELECT * FROM
Instructor
WHERE fName LIKE 'J%';
- (14) SELECT * FROM
Instructor
WHERE fName LIKE 'J_h_';
- (15) SELECT fName, lName FROM
Instructor
WHERE address LIKE '%Houston,TX%';
- (16) SELECT (salary+bonus) FROM
Instructor;
- (17) SELECT (salary+bonus) AS
total_income FROM Instructor;
- (18) SELECT ((A1/3.14)*2.54) AS
A1_IN_CM FROM R;
- (19) SELECT SYSDATE AS TODAY
FROM DUAL;
- (20) SELECT FName, Lname, Bday,
TRUNC (MONTHS_BETWEEN (SYSDATE, Bday)/12) AS "Actual Age"
FROM Person;

DATE Manipulation

- Show today's date in the form of MM/DD/YYYY format
SELECT SYSDATE, TO_CHAR(SYSDATE, 'MM/DD/YYYY') as
CurrDate FROM DUAL;
- Show today's date including time
SELECT SYSDATE, TO_CHAR(SYSDATE, 'yyyy-mm-dd hh24:mi:ss')
as CurrDateTime FROM DUAL;
- Find products whose INDATE is more than 90 days old
SELECT P_Code, P_Desc, P_Indate
FROM Product
WHERE P_Indate <= SYSDATE - 90;
- Find all the orders received in Feb 2012.
SELECT Order#, Odate
FROM ORDER
WHERE Odate BETWEEN '01-Feb-12' AND '29-Feb-12';

STRING Manipulation

- Concatenation**

```
SQL> SELECT Lname||', '||Fname "FULL NAME" FROM Emp;
will display
FULL NAME
Bond, James
```

- Capitalization**

```
SQL> SELECT Lname, INITCAP(Lname), UPPER(Lname), LOWER(Lname)
FROM Emp;
```

- Capitalization:** Show only the first letter in capital and the rest in lower

```
SQL> SELECT INITCAP(LOWER(Lname)) FROM Emp;
```

/* Padding blanks to strings*/

RPAD: Pad to the right side of the column with left justification

LPAD: Pad to the left side of the column with right justification

```
SQL> SELECT RPAD(Lname, 15, '.'), Age, LPAD(Lname, 15), FROM EMP:
Clinton..... 53 .....Clinton
Gore..... 52 .....Gore
```

STRING Manipulation

/* **TRIM** function is used to remove all leading or trailing characters (or both) from a character string.*/

```
SQL>SELECT TRIM(" LEADS Fellow ") FROM DUAL;
```

/* Use of **LTRIM**(left trim) and **RTRIM**(right trim) functions*/

```
SQL> SELECT Lname, LTRIM(Lname, 'SA'), RTRIM(Lname, 'S')
FROM Emp;
(LTRIM prints M from SAM, RTRIM prints WALE from WALES)
```

/* Remove "." at the end and "THE from the front */

```
SQL> SELECT LTRIM ( RTRIM (Title, '.'), 'THE ') FROM Magazine;
```

Will convert

MY DARING."

"THE GOD FATHER."

into

MY DARING

GOD FATHER

STRING Manipulation

/* INSTR(string, set [, start [, occurrence]]) tells you where in the string it found what you were searching for */

/*Show names and position which have 'E' in the 5th or greater position */

```
SQL> SELECT Lname, INSTR(Lname, 'E', 5) "AFTER FIVE"
      FROM Emp;
```

/* SUBSTR(String, starting_position, #chars to be returned)

Extract numeric part, add 1000, and concatenate */

```
SQL> SELECT Lname, E_ID,
      'S' || TO_CHAR( TO_NUMBER (SUBSTR (E_ID,2,3) ) + 1000) "NEW E_#"
      FROM Emp;
```

Will display

LNAME	E_ID	NEW E_#
Jones	E001	S1001
Wales	E002	S1002

Aggregation

- Example of five aggregation functions

```
SELECT MAX(bonus), MIN(bonus), AVG(bonus), COUNT(bonus),
SUM(bonus)
FROM Instructor
WHERE position = 'assistant';
```

- Note: What's wrong with the following?

```
SELECT instructorID, bonus
FROM Instructor
WHERE bonus > AVG(bonus);
```

We can use aggregate functions only in **SELECT** and **HAVING** clause

- How many different course titles are there?

```
SELECT COUNT(DISTINCT title) AS count
FROM Course;
```

Descriptive Statistics in SQL

- **SELECT** cus_fname, cus_lname, cus_balance
AVG(cus_balance) mean,
MEDIAN (cus_balance) median,
STATS_MODE (cus_balance) mode,
STDDEV (cus_balance) “Standard Deviation”
FROM Product;

CUSTOMER	
PK	<u>cus_code</u>
	cus_lname cus_fname cus_initial cus_areacode cus_phone cus_balance

- Computing median over partition
SELECT cus_fname, cus_lname, cus_balance
MEDIAN (cus_balance) **OVER** (**PARTITION BY** cus_areacode) AS median
FROM Product;

- Computing SD over unique values
SELECT cus_fname, cus_lname, cus_balance
STDDEV (**DISTINCT** cus_balance) **FROM** Product;

Aggregation with GROUP BY

Find the total number of instructors in each department and the sum of their bonus, respectively

```
SELECT    deptCode, COUNT(instructorID) AS count, SUM(bonus) AS sum
FROM      Instructor
GROUP BY  deptCode
ORDER BY  deptCode;
```

- Result:

deptCode	count	sum
acct	2	1100
math	2	300

Instructor
-instructorID(PK)
-fName
-lName
-ssn
-deptCode
-deptName
-position
-bonus

- GROUP BY** must include *all non-aggregate function column names in the SELECT list.*

Aggregation with GROUP BY and HAVING

```

Oracle SQL*Plus
File Edit Search Options Help
SQL> SELECT U_CODE, COUNT(DISTINCT (P_CODE)), AVG(P_PRICE)
2 FROM PRODUCT
3 GROUP BY U_CODE;

U_CODE COUNT(DISTINCT(P_CODE)) AVG(P_PRICE)
-----
21225      2      8.47
21231      1      8.45
21344      3     12.49
23119      2     41.97
24288      3    155.593333
25595      3      89.63
          2     10.135

7 rows selected.

SQL> SELECT U_CODE, COUNT(DISTINCT (P_CODE)), AVG(P_PRICE)
2 FROM PRODUCT
3 GROUP BY U_CODE
4 HAVING AVG(P_PRICE) < 10;

U_CODE COUNT(DISTINCT(P_CODE)) AVG(P_PRICE)
-----
21225      2      8.47
21231      1      8.45

```

PRODUCT	
PK	<u>p_code</u>
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

Aggregation with GROUP BY and HAVING

- Find the number of invoice per year and per month for the last 3 years

```

SELECT EXTRACT(year FROM inv_date) "Year",
       EXTRACT(month FROM inv_date) "Month",
       COUNT(inv_date) "No. of Invoices"
FROM invoice
GROUP BY EXTRACT(year FROM inv_date),
         EXTRACT(month FROM inv_date)
HAVING EXTRACT(year FROM inv_date) IN (2017, 2016, 2015)
ORDER BY "No. of Invoices" DESC;

```

INVOICE	
PK	<u>inv_number</u>
FK1	cus_code
	inv_date

CASE Statement

- Simple CASE Statements

```
SELECT      Fname, Lname,
            (CASE DNO
              WHEN 1      THEN 'Headquarters'
              WHEN 4      THEN 'Administration'
              WHEN 5      THEN 'Research'
              ELSE        'No department'
            END) AS Department
FROM      Employee;
```

Output:

<u>Fname</u>	<u>Lname</u>	<u>Department</u>
John	Smith	Research
Franklin	Wong	Research
Alica	Zelaya	Administration

CASE Statement

- Searched CASE Statements

```
SELECT      Fname, Lname, Salary
            (CASE Salary
              WHEN Salary <= 25000 THEN 1500
              WHEN Salary > 25000 AND Salary < 50000 THEN 1000
              WHEN Salary > 50000 AND Salary < 100000 THEN 500
              ELSE 0
            END) "Bonus"
FROM      Employee;
```

Output:

<u>Fname</u>	<u>Lname</u>	<u>Salary</u>	<u>Bonus</u>
John	Smith	30000	1000
Franklin	Wong	40000	1000
Alica	Zelaya	25000	1500

Nested Subquery and Aggregate Functions

- Find the product that has the oldest date

```
SELECT P_CODE, P_DESCRIPT
FROM PRODUCT
WHERE P_INDATE = (
    SELECT MIN (P_INDATE)
    FROM PRODUCT);
```

PRODUCT	
PK	<u>p_code</u>
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code


This query is in the form of a nested query.

JOIN: Two Syntax

JOIN is usually done through **PK-FK chains**

- For each product, find their vendor code and names

```
SELECT      P.P_Code, V.V_Code, V.V_Name
FROM        Product P, Vendor V
WHERE       P.V_Code = V.V_code;
```



```
SELECT      P.P_Code, V.V_Code, V.V_Name
FROM        Product P INNER JOIN Vendor V
ON          P.V_Code = V.V_code;
```


Types of JOINS

TABLE 8.1


SQL JOIN EXPRESSION STYLES

JOIN CLASSIFICATION	JOIN TYPE	SQL SYNTAX EXAMPLE	DESCRIPTION
CROSS	CROSS JOIN	SELECT * FROM T1, T2	Returns the Cartesian product of T1 and T2 (old style)
		SELECT * FROM T1 CROSS JOIN T2	Returns the Cartesian product of T1 and T2
INNER	Old-style JOIN	SELECT * FROM T1, T2 WHERE T1.C1=T2.C1	Returns only the rows that meet the join condition in the WHERE clause (old style); only rows with matching values are selected
	NATURAL JOIN	SELECT * FROM T1 NATURAL JOIN T2	Returns only the rows with matching values in the matching columns; the matching columns must have the same names and similar data types
	JOIN USING	SELECT * FROM T1 JOIN T2 USING (C1)	Returns only the rows with matching values in the columns indicated in the USING clause
	JOIN ON	SELECT * FROM T1 JOIN T2 ON T1.C1=T2.C1	Returns only the rows that meet the join condition indicated in the ON clause
OUTER	LEFT JOIN	SELECT * FROM T1 LEFT OUTER JOIN T2 ON T1.C1=T2.C1	Returns rows with matching values and includes all rows from the left table (T1) with unmatched values
	RIGHT JOIN	SELECT * FROM T1 RIGHT OUTER JOIN T2 ON T1.C1=T2.C1	Returns rows with matching values and includes all rows from the right table (T2) with unmatched values
	FULL JOIN	SELECT * FROM T1 FULL OUTER JOIN T2 ON T1.C1=T2.C1	Returns rows with matching values and includes all rows from both tables (T1 and T2) with unmatched values


Visualizing JOIN Operations




SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key




SELECT <auswahl>
FROM tabelleA A
INNER JOIN tabelleB B
ON A.key = B.key





SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key



SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
WHERE B.key IS NULL




SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key



SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL

INNER JOIN

- For each invoice, find the product description, #units, and price of items in the invoice.



SQL Plus

```

SQL> SELECT INVOICE.INV_NUMBER, PRODUCT.P_CODE, P_DESCRIPT, LINE_UNITS, LINE_PRICE
2 FROM INVOICE JOIN LINE ON INVOICE.INV_NUMBER = LINE.INV_NUMBER
3 JOIN PRODUCT ON LINE.P_CODE = PRODUCT.P_CODE;
  
```


INV_NUMBER	P_CODE	P_DESCRIPT	LINE_UNITS	LINE_PRICE
1001	13-Q2/P2	7.25-in. pwr. saw blade	1	14.99
1001	23109-HB	Claw hammer	1	9.95
1002	54778-2T	Rat-tail file, 1/8-in. fine	2	4.99
1003	2238/QPD	B&D cordless drill, 1/2-in.	1	38.95
1003	1546-QQ2	Hrd. cloth, 1/4-in., 2x50	1	39.95
1003	13-Q2/P2	7.25-in. pwr. saw blade	5	14.99
1004	54778-2T	Rat-tail file, 1/8-in. fine	3	4.99
1004	23109-HB	Claw hammer	2	9.95
1005	PVC23DRT	PVC pipe, 3.5-in., 8-ft	12	5.87
1006	SM-18277	1.25-in. metal screw, 25	3	6.99
1006	2232/QTY	B&D jigsaw, 12-in. blade	1	109.92
1006	23109-HB	Claw hammer	1	9.95
1006	89-WRE-Q	Hicut chain saw, 16 in.	1	256.99
1007	13-Q2/P2	7.25-in. pwr. saw blade	2	14.99
1007	54778-2T	Rat-tail file, 1/8-in. fine	1	4.99
1008	PVC23DRT	PVC pipe, 3.5-in., 8-ft	5	5.87
1008	WR3/TT3	Steel matting, 4'x8'x1/6", .5" mesh	3	119.95
1008	23109-HB	Claw hammer	1	9.95

18 rows selected.

SQL>

LEFT OUTER JOIN

- For ALL the Vendors, find all the products they provide. Include all the vendors whether or not they provide any product.



SQL Plus

```

SQL> SELECT P_CODE, VENDOR.V_CODE, V_NAME
2 FROM VENDOR LEFT JOIN PRODUCT ON VENDOR.V_CODE = PRODUCT.V_CODE;
  
```


P_CODE	V_CODE	V_NAME
110ER/31	25595	Rubicon Systems
13-Q2/P2	21344	Gomez Bros.
14-Q1/L3	21344	Gomez Bros.
1546-QQ2	23119	Randssets Ltd.
1558-QW1	23119	Randssets Ltd.
2232/QTY	24288	ORDVA, Inc.
2232/QWE	24288	ORDVA, Inc.
2238/QPD	25595	Rubicon Systems
23109-HB	21225	Bryson, Inc.
54778-2T	21344	Gomez Bros.
89-WRE-Q	24288	ORDVA, Inc.
SM-18277	21225	Bryson, Inc.
SW-23116	21231	D&E Supply
WR3/TT3	25595	Rubicon Systems
	22567	Dome Supply
	21226	SuperLo0, Inc.
	24004	Brackman Bros.
	25501	Damal Supplies
	25443	B&K, Inc.

19 rows selected.

SQL>

SET Operations


```
-- All customers and staff
SELECT first_name, last_name
FROM customer
UNION
SELECT first_name, last_name
FROM staff
```



UNION

 $R \cup S$


```
-- customers that are actors
SELECT first_name, last_name
FROM customer
INTERSECT
SELECT first_name, last_name
FROM actor
```



INTERSECT

 $R \cap S$

```
-- customers that are not staff
SELECT first_name, last_name
FROM customer
EXCEPT
SELECT first_name, last_name
FROM staff
```



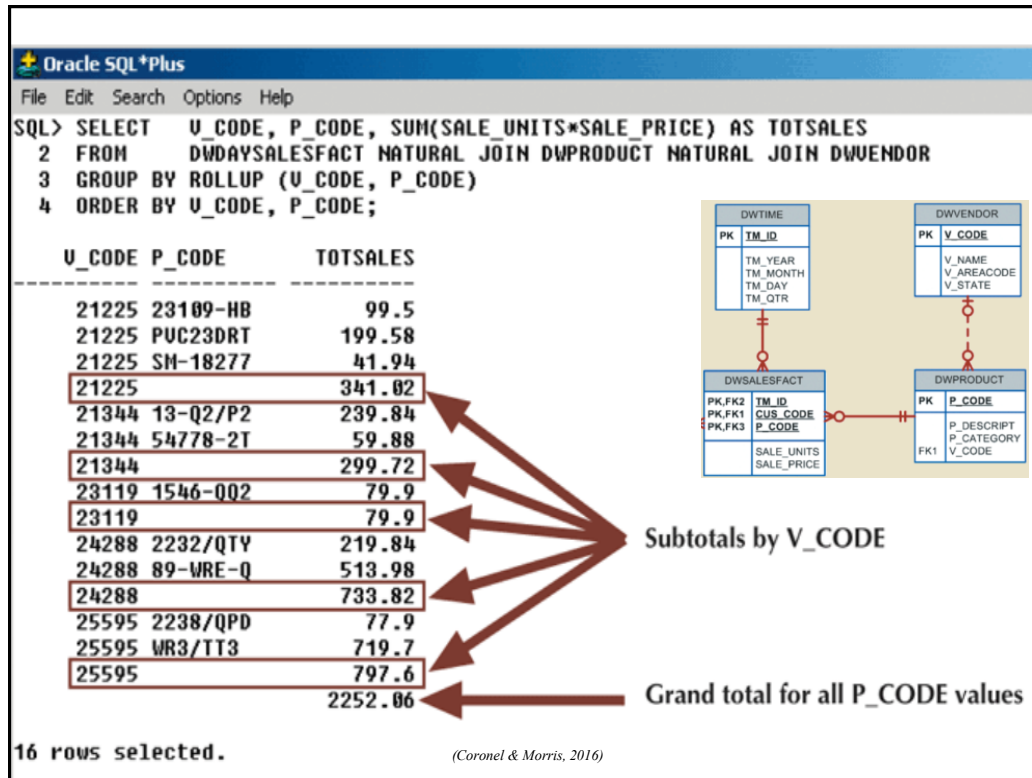
MINUS

 $R - S$

SQL Analytic Functions: ROLLUP

- Example: **GROUP BY ROLLUP** (V_Code, P_Code) produces the union of
 - GROUP BY V_Code, P_Code
 - GROUP BY V_Code
 - Grand total

```
SQL> SELECT  V_CODE, P_CODE, SUM(SALE_UNITS*SALE_PRICE) AS TOTSALES
2 FROM      DWDALESFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWVENDOR
3 GROUP BY  ROLLUP (V_CODE, P_CODE)
4 ORDER BY  V_CODE, P_CODE;
```



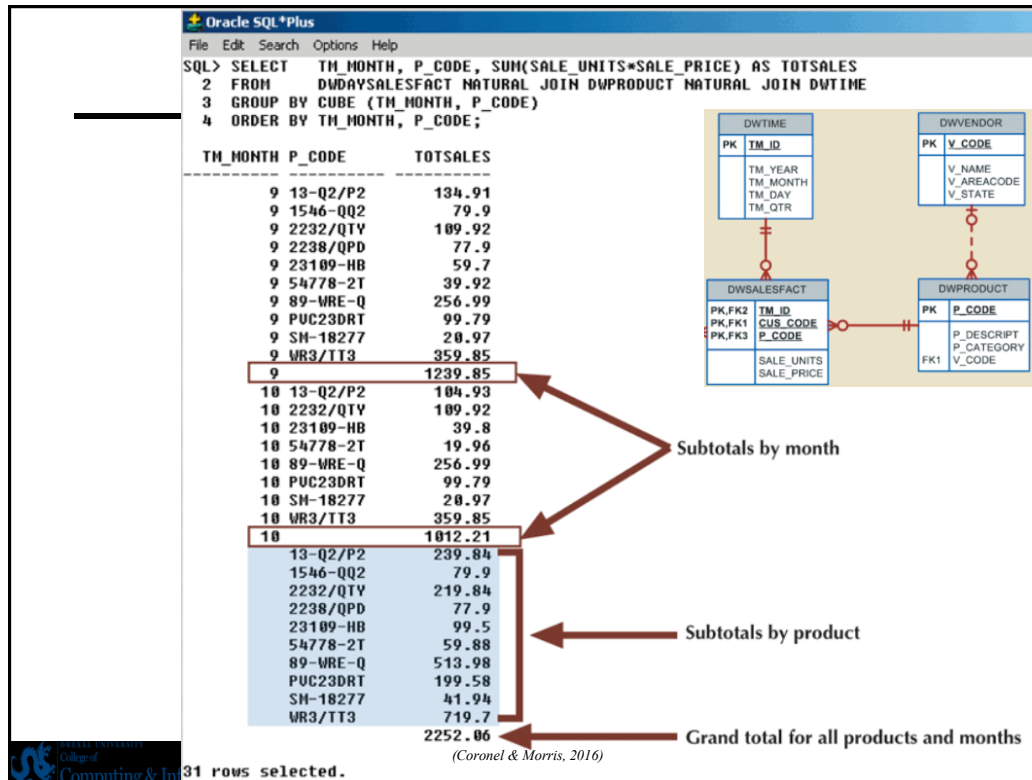
SQL Analytic Functions: Cube

- The CUBE extension
 - Enables you to get a subtotal for each column listed in the expression, in addition to a grand total for the last column listed
 - GROUP BY CUBE (TM_Month, P_CODE) is a union of:
 - GROUP BY CUBE (TM_Month, P_CODE)
 - GROUP BY CUBE (TM_Month)
 - GROUP BY CUBE (P_CODE)
 - Grand total

```

SQL> SELECT  TM_MONTH, P_CODE, SUM(SALE_UNITS*SALE_PRICE) AS TOTSALES
2 FROM      DWDAYSALSALESFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWTIME
3 GROUP BY CUBE (TM_MONTH, P_CODE)
4 ORDER BY TM_MONTH, P_CODE;

```



RANK and DENSE RANK Functions

- RANK() leave ranking gaps when there multiple rows in the same rank (1, 2, 2, 4, 5...)
- DENSE_RANK() does **not** leave ranking gaps (1, 2, 2, 3, 4, 5)

```

SELECT p_code, p_descript, p_price,
       RANK() OVER
         (ORDER BY p_price NULLS LAST) AS Rank,
       DENSE_RANK ()
         (ORDER BY p_price NULLS LAST) AS Dense_rank
FROM   Product
ORDER BY p_price;

```

PRODUCT	
PK	<u>p_code</u>
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

- DESC means descending order
- NULLS LAST means null values are smaller than non-null values
- You may use NULLS FIRST

RANK and DENSE RANK Functions

- Ranking applied to dates

```
SELECT p_code, p_descript, p_indate,
       RANK() OVER
         (ORDER BY p_indate NULLS LAST) AS Rank,
       DENSE_RANK ()
         (ORDER BY p_indate NULLS LAST) AS Dense_rank
FROM   Product
ORDER BY p_indate;
```

PRODUCT	
PK	<u>p_code</u>
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

Top-N Queries

- Before Oracle 12c:

```
SELECT p_code, p_descript,
       (SELECT p_price,
        RANK() OVER
          (ORDER BY p_price NULLS LAST) AS Rank
        FROM product)
FROM   Product
WHERE  Rank <= 10
ORDER BY p_price;
```

- In Oracle 12c:

```
SELECT p_code, p_descript, p_price,
       RANK() OVER
         (ORDER BY p_price NULLS LAST) AS Rank,
FROM   Product
ORDER BY p_price;
FETCH FIRST 10 ROWS ONLY;
```

- Another option by percentage
FETCH FIRST 20 PERCENT ROWS ONLY;
- MySQL uses LIMIT clause
LIMIT 10

Question?

